

TVM: A Tile-based Video Management Framework


Tianxiong Zhong¹, Zhiwei Zhang¹, Guo Lu², Ye Yuan¹, Yu-Ping Wang¹, and Guoren Wang¹

¹Beijing Institute of Technology ²Shanghai JiaoTong University



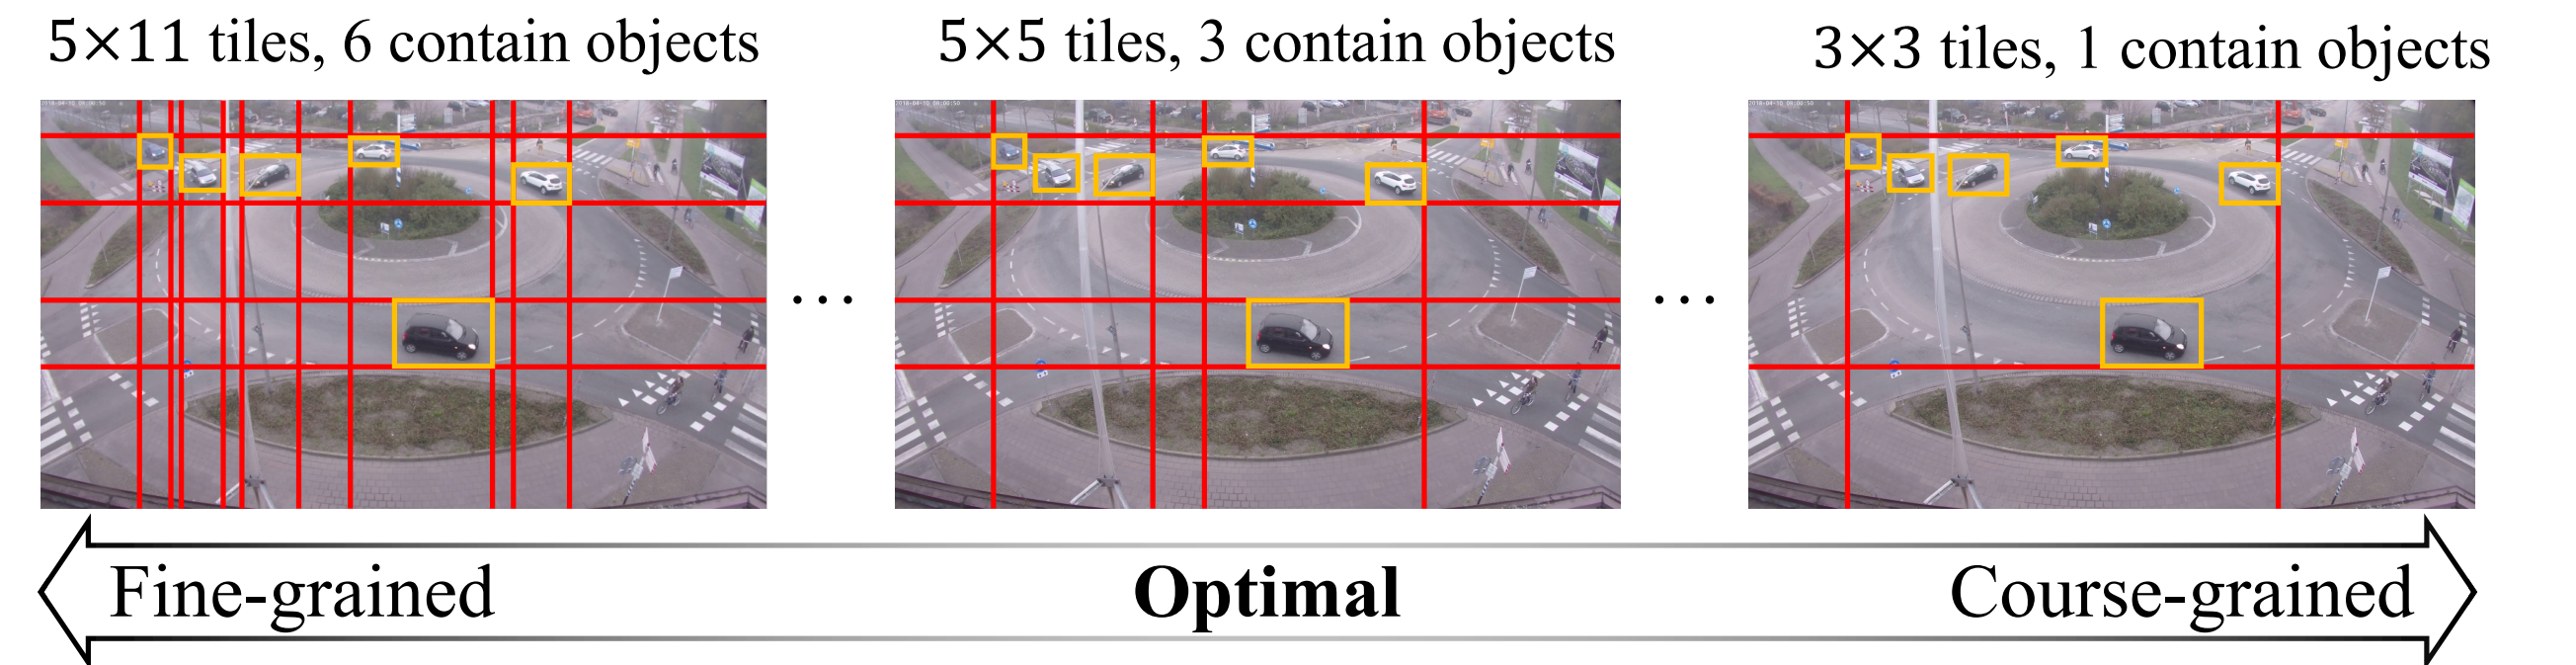
Motivation

Query: How many cars are there in the video? Decode+DNN



By reducing data access from the spatial scale to accelerate video analysis queries.


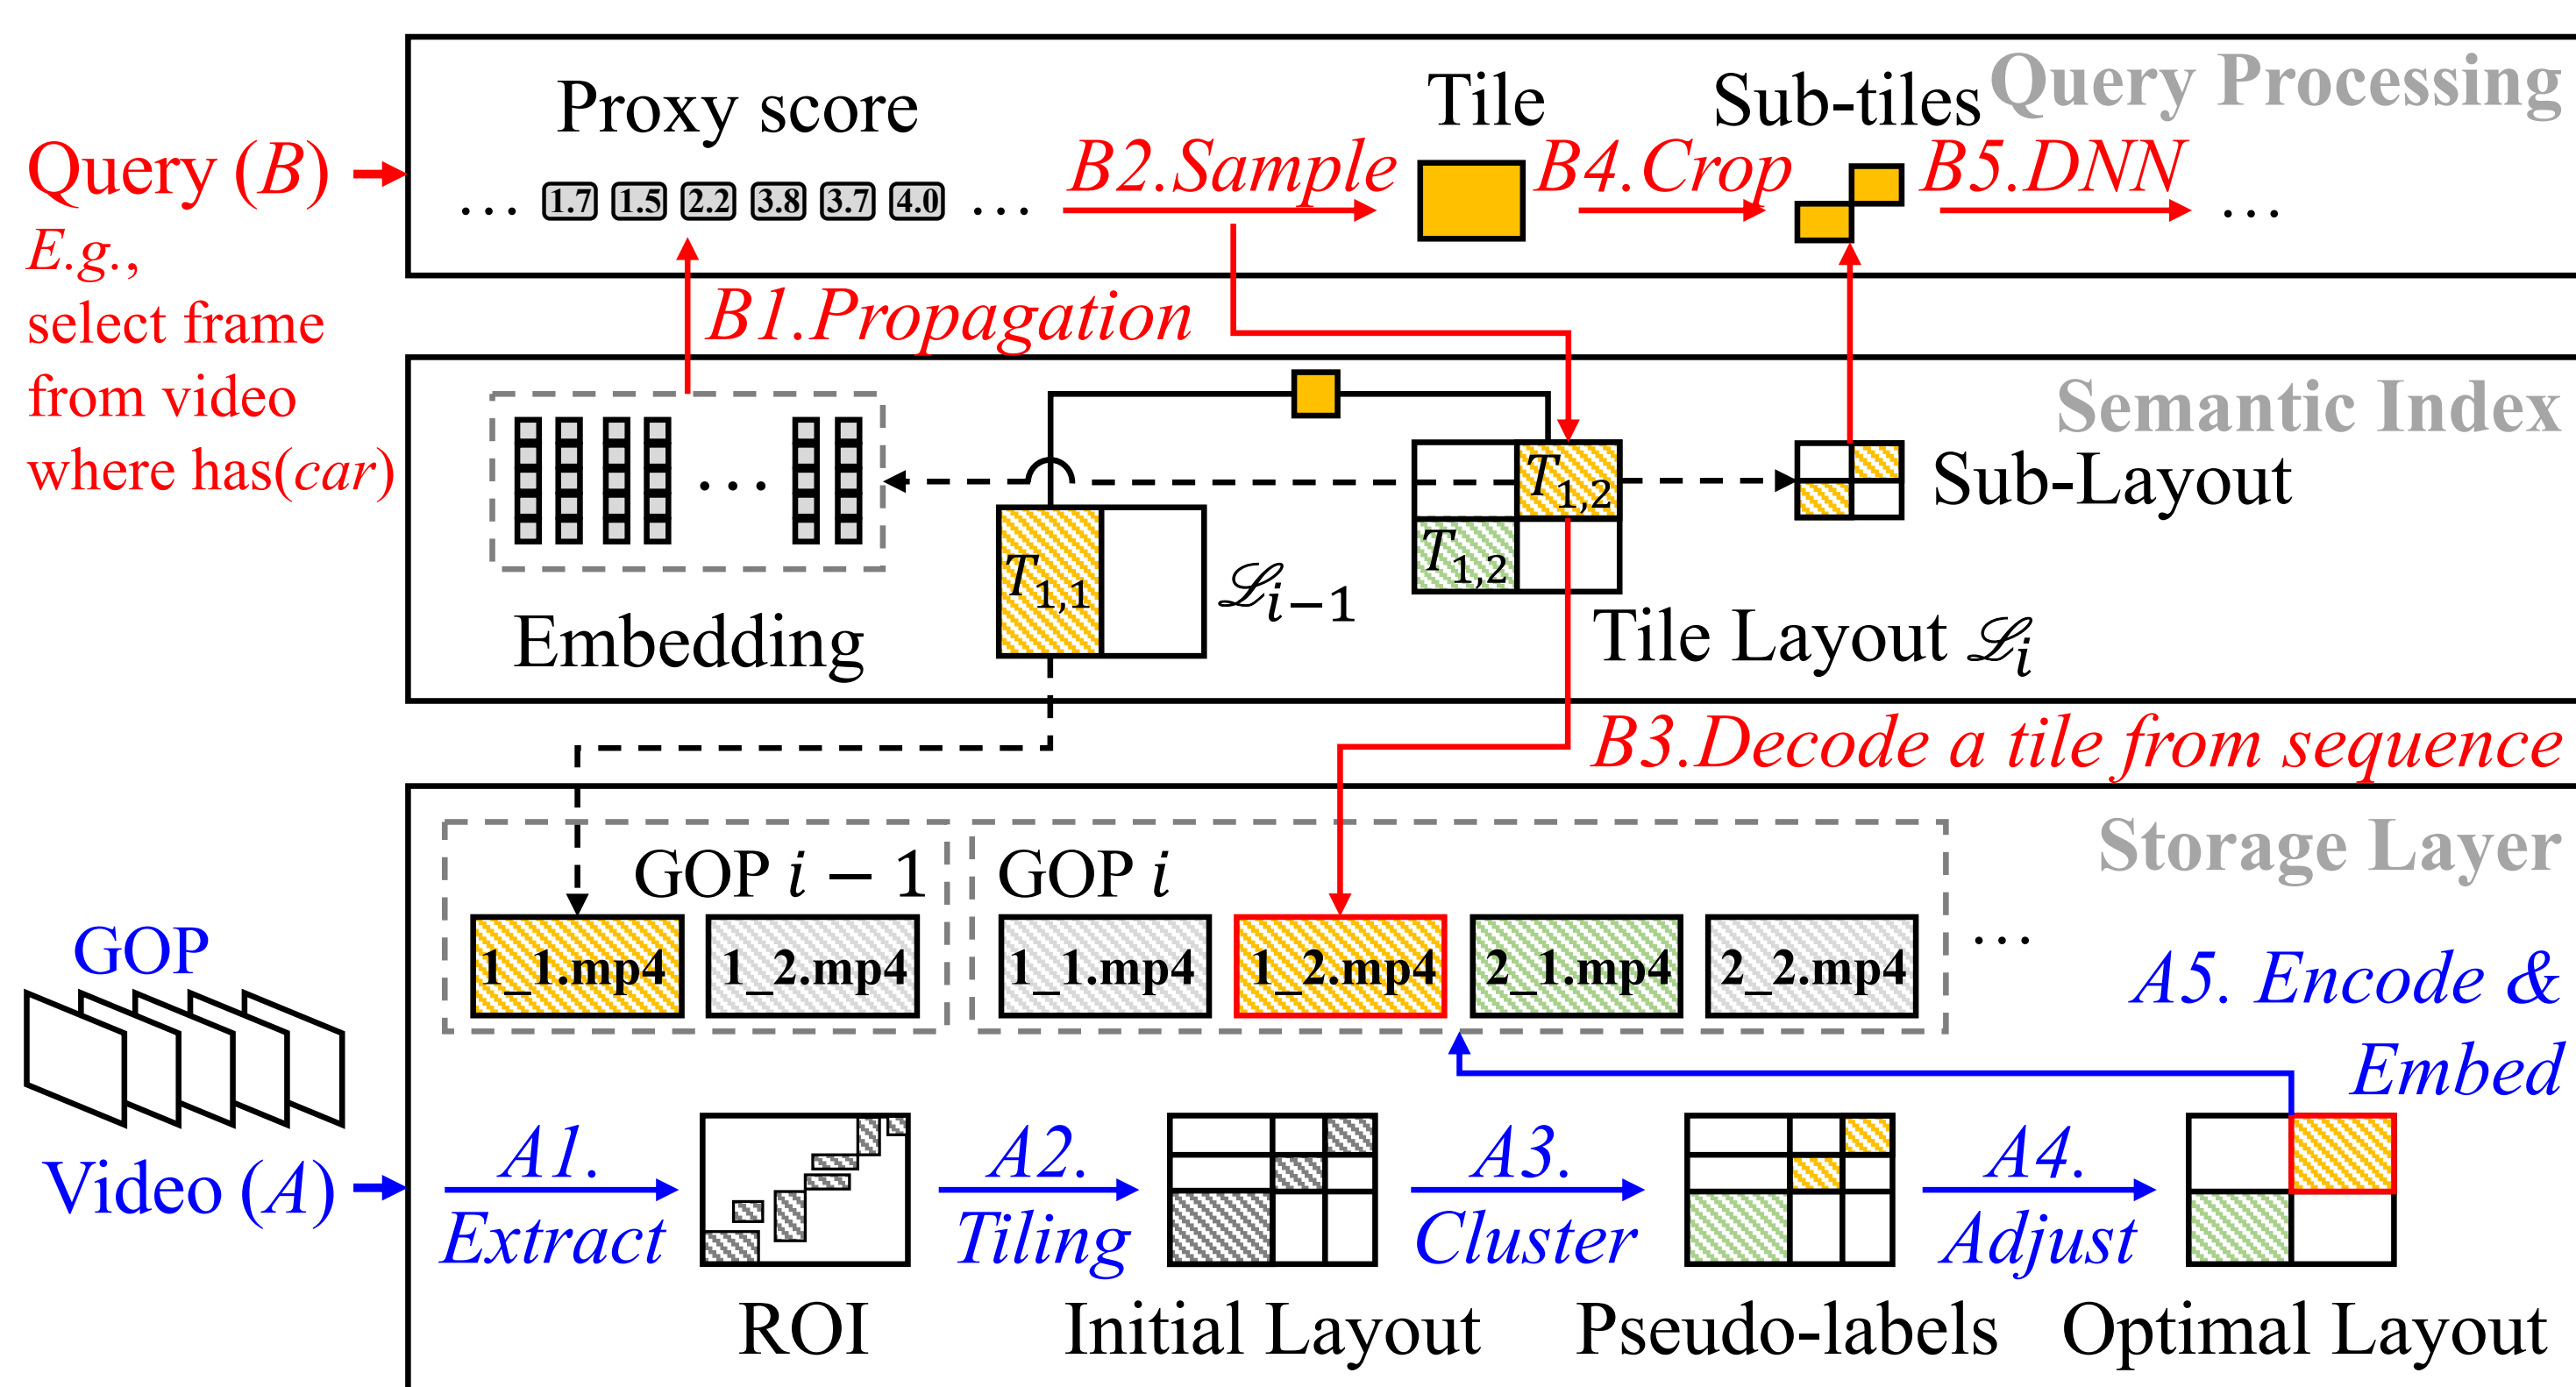
Challenge



- ✓ Small tile size
- ✓ Suitable tile size
- ✗ Huge tile size
- ✗ Decode many tiles
- ✓ Decode less tiles
- ✓ Decode only one tile

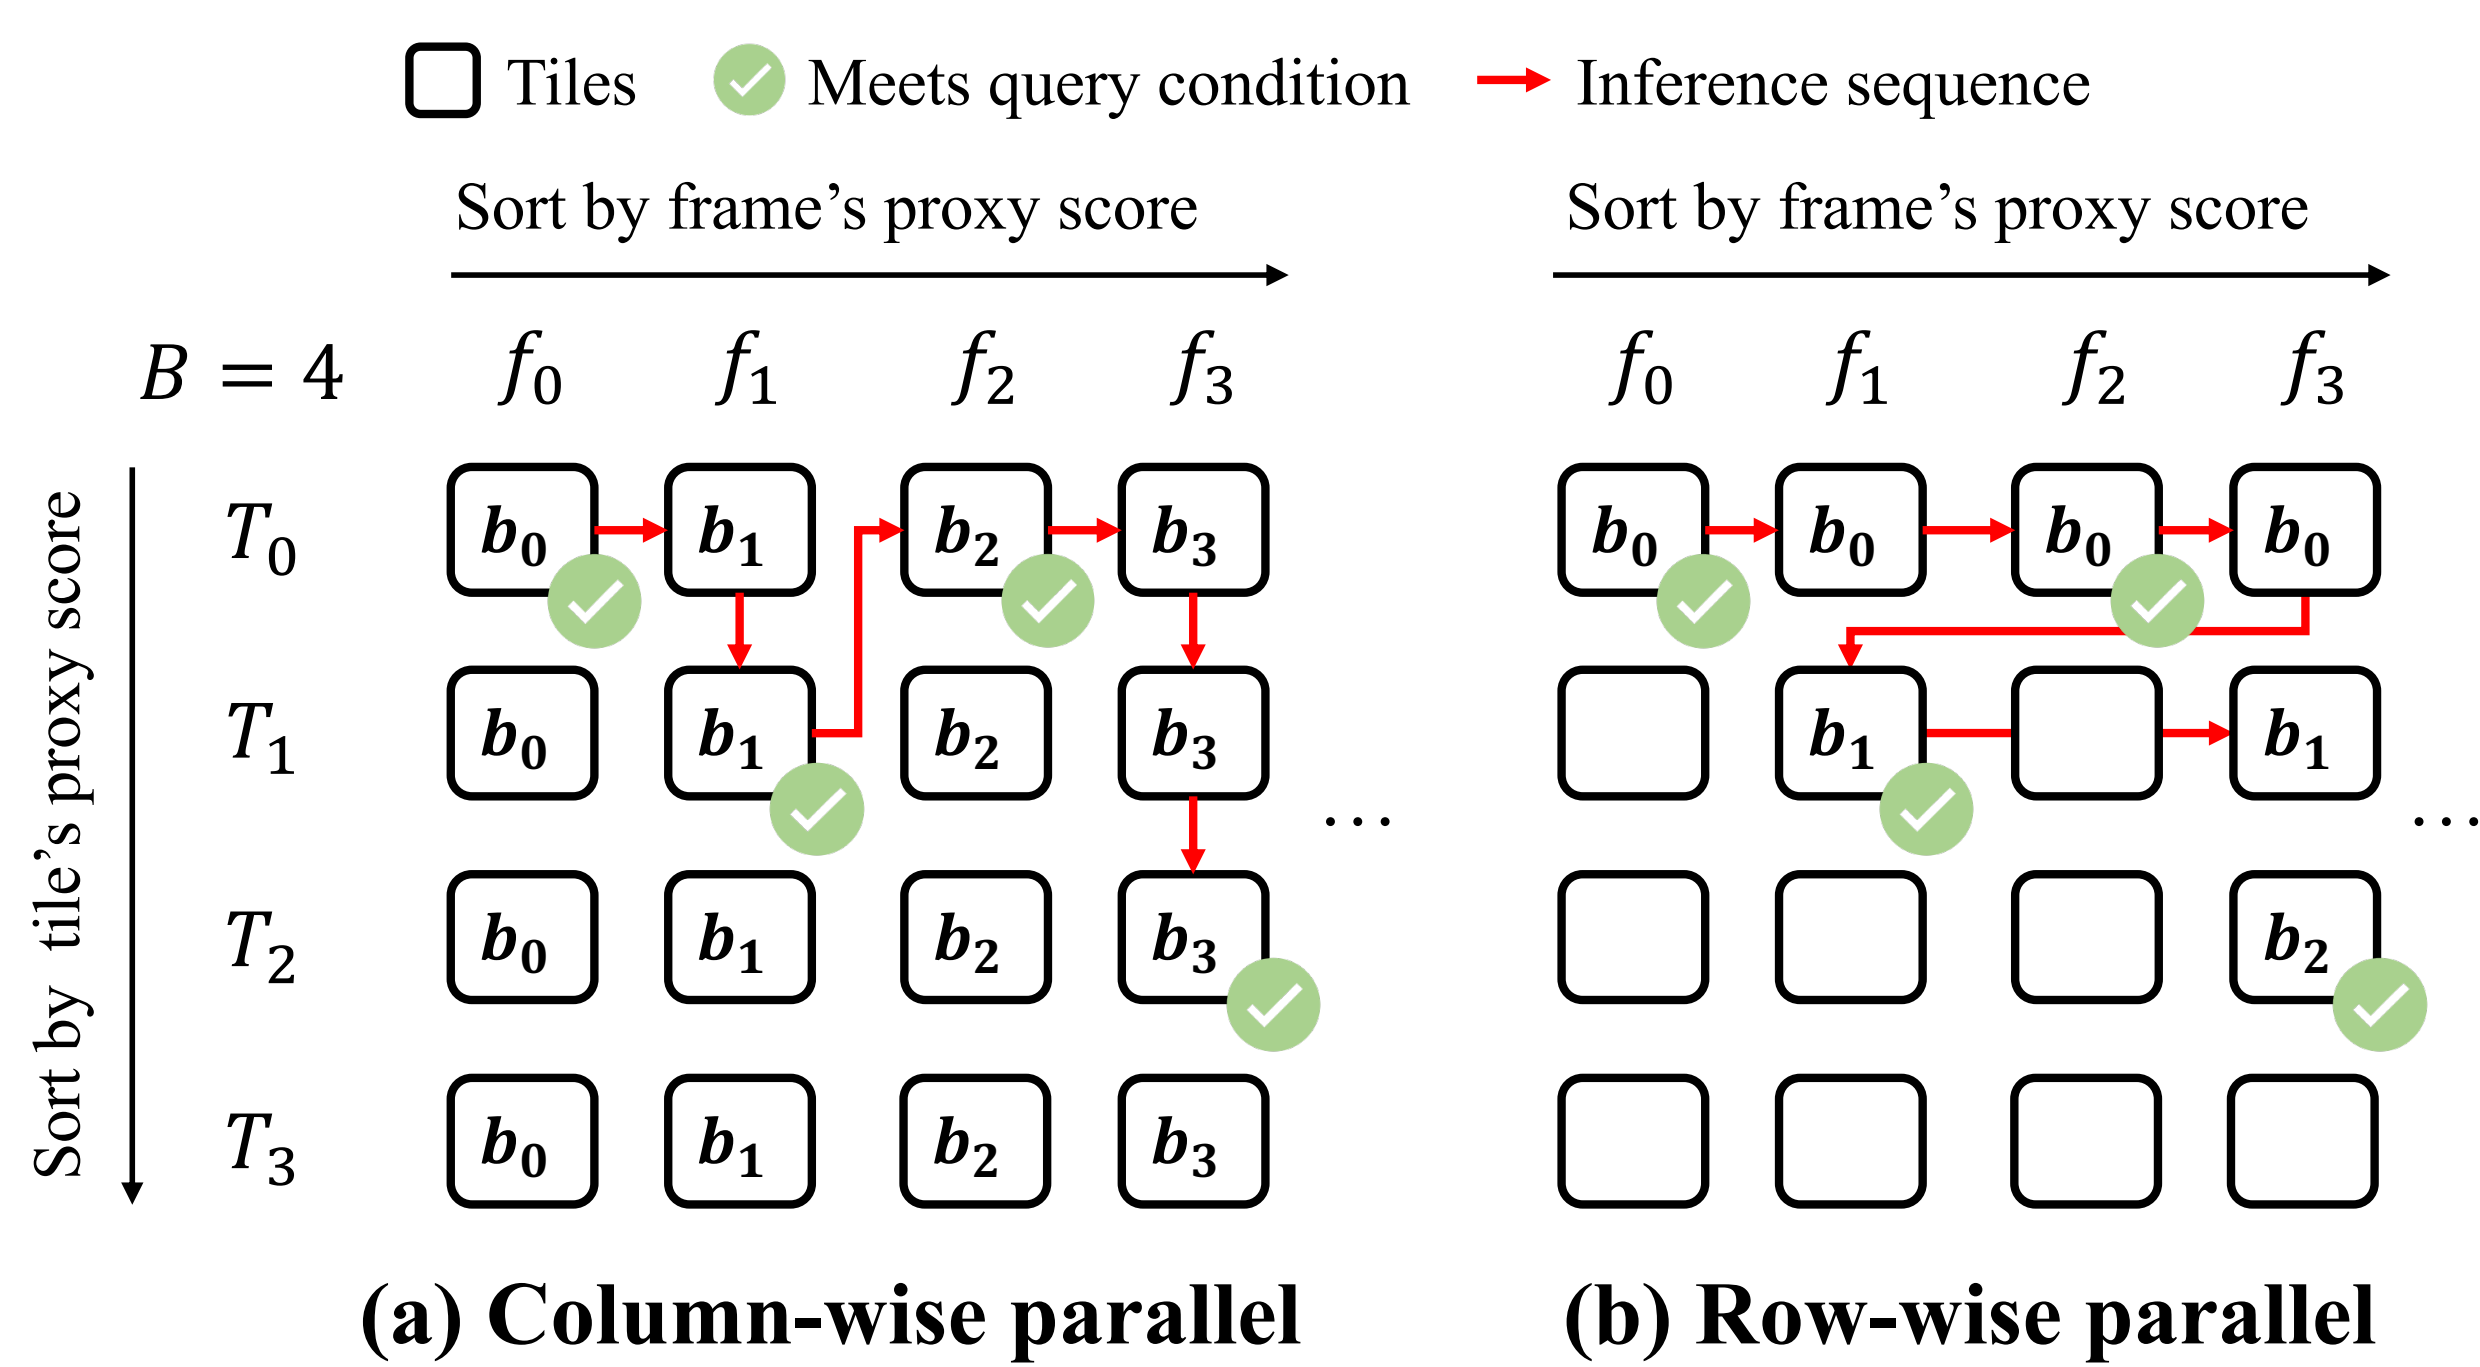
- Design of an optimal tile layout that minimizes the overall decoding and processing costs.
- Query workflows for accessing tiles are different. Need query-specific acceleration algorithms.

Optimal Index Structure




- We model the decoding/DNN cost as a function of the number of pixels which can be used to quantify the pros and cons of different layouts.
- To speed up the search for the optimal layout, we also propose DP and greedy algorithms.

Query-Specific Acceleration Algorithms



Parallel decoding algorithm for Limit queries to avoid access tiles that do not need to be processed.



Cache decoders, decoding states, and dependent frames to alleviate the resource waste in track queries.

Experiment

